

Muen – Design

Inzemamul Haque

25 Nov 2016

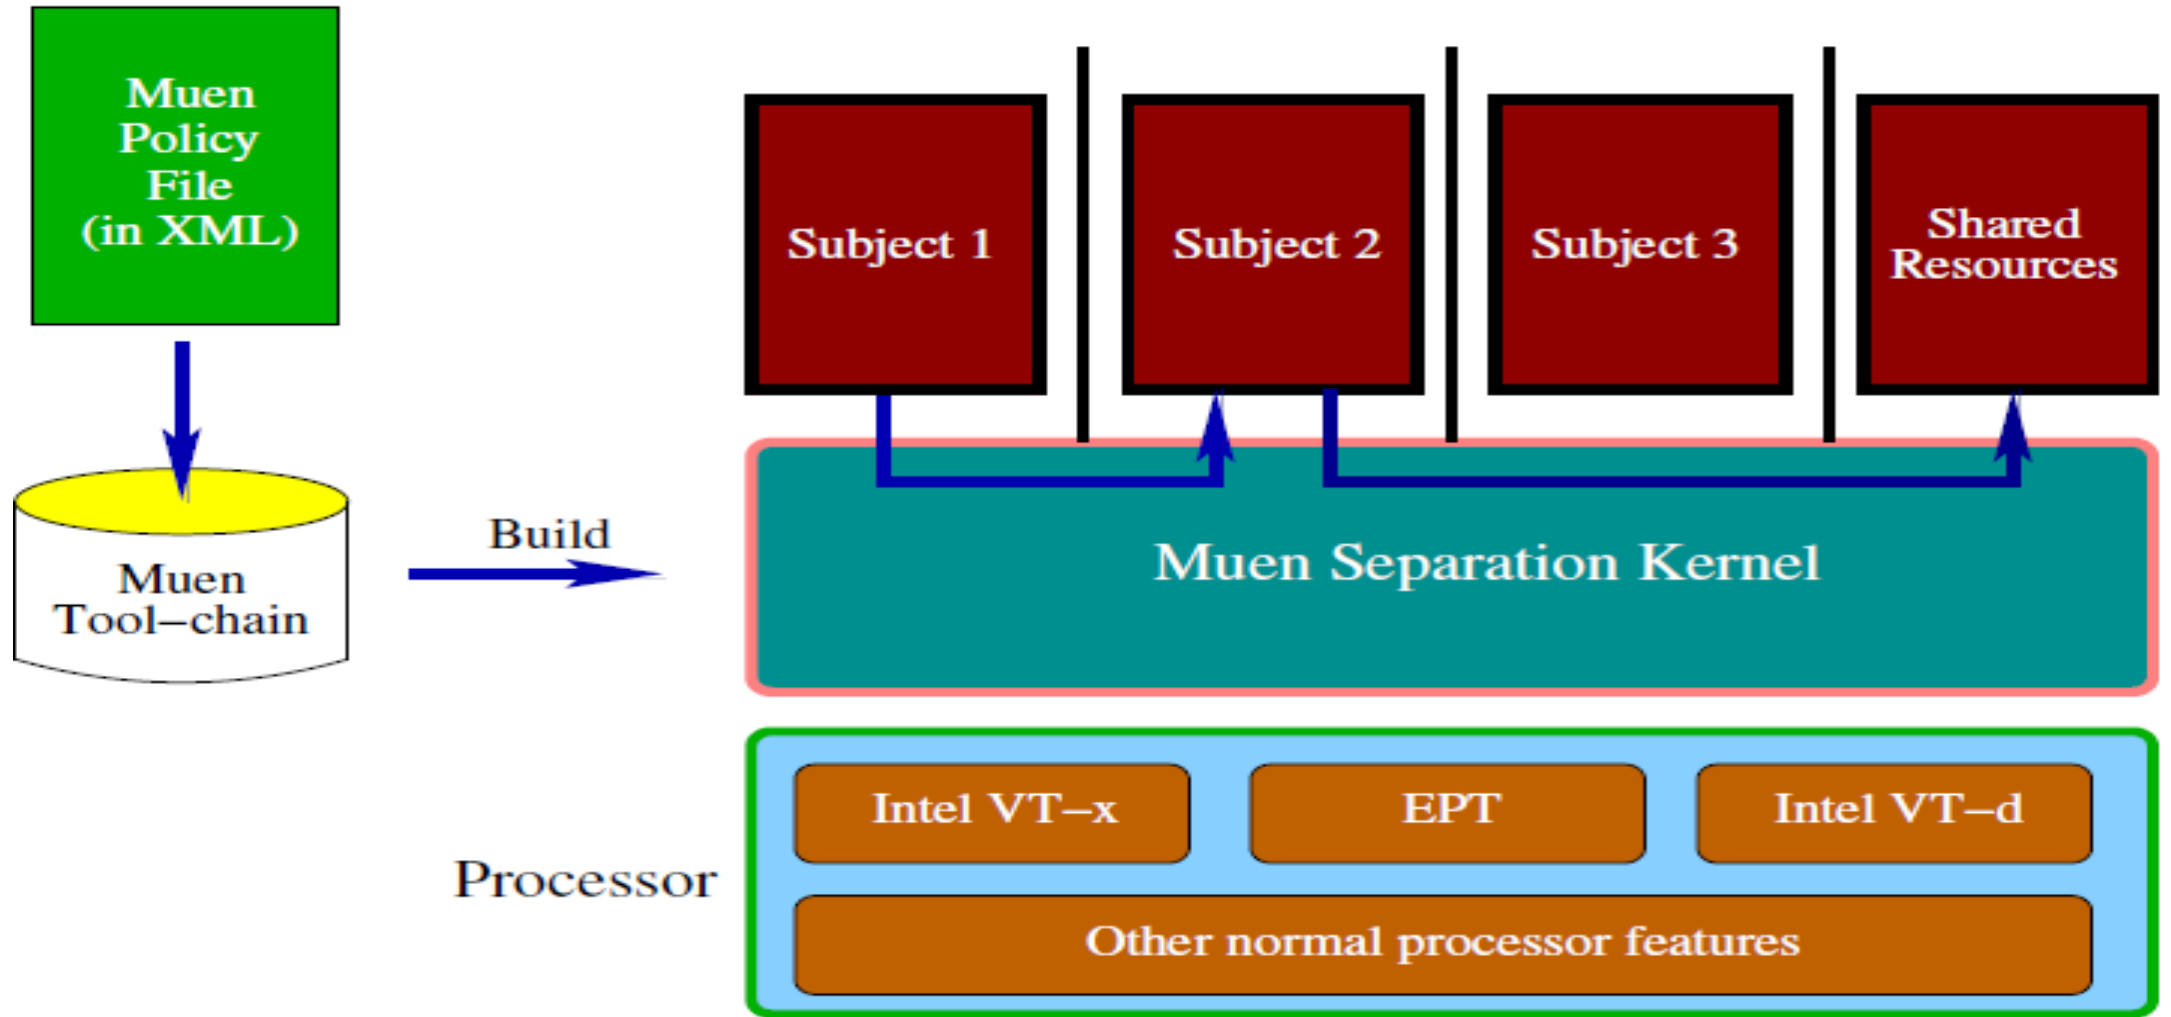
Introduction

- Muen is an open-source separation kernel for x86 platform
- Uses Intel hardware support for virtualization

What Muen does?

- Takes a policy as input and works according to it
- Policy contains information like
 - No. of subjects
 - Information about memory
 - Scheduling policy
 - Communication channels

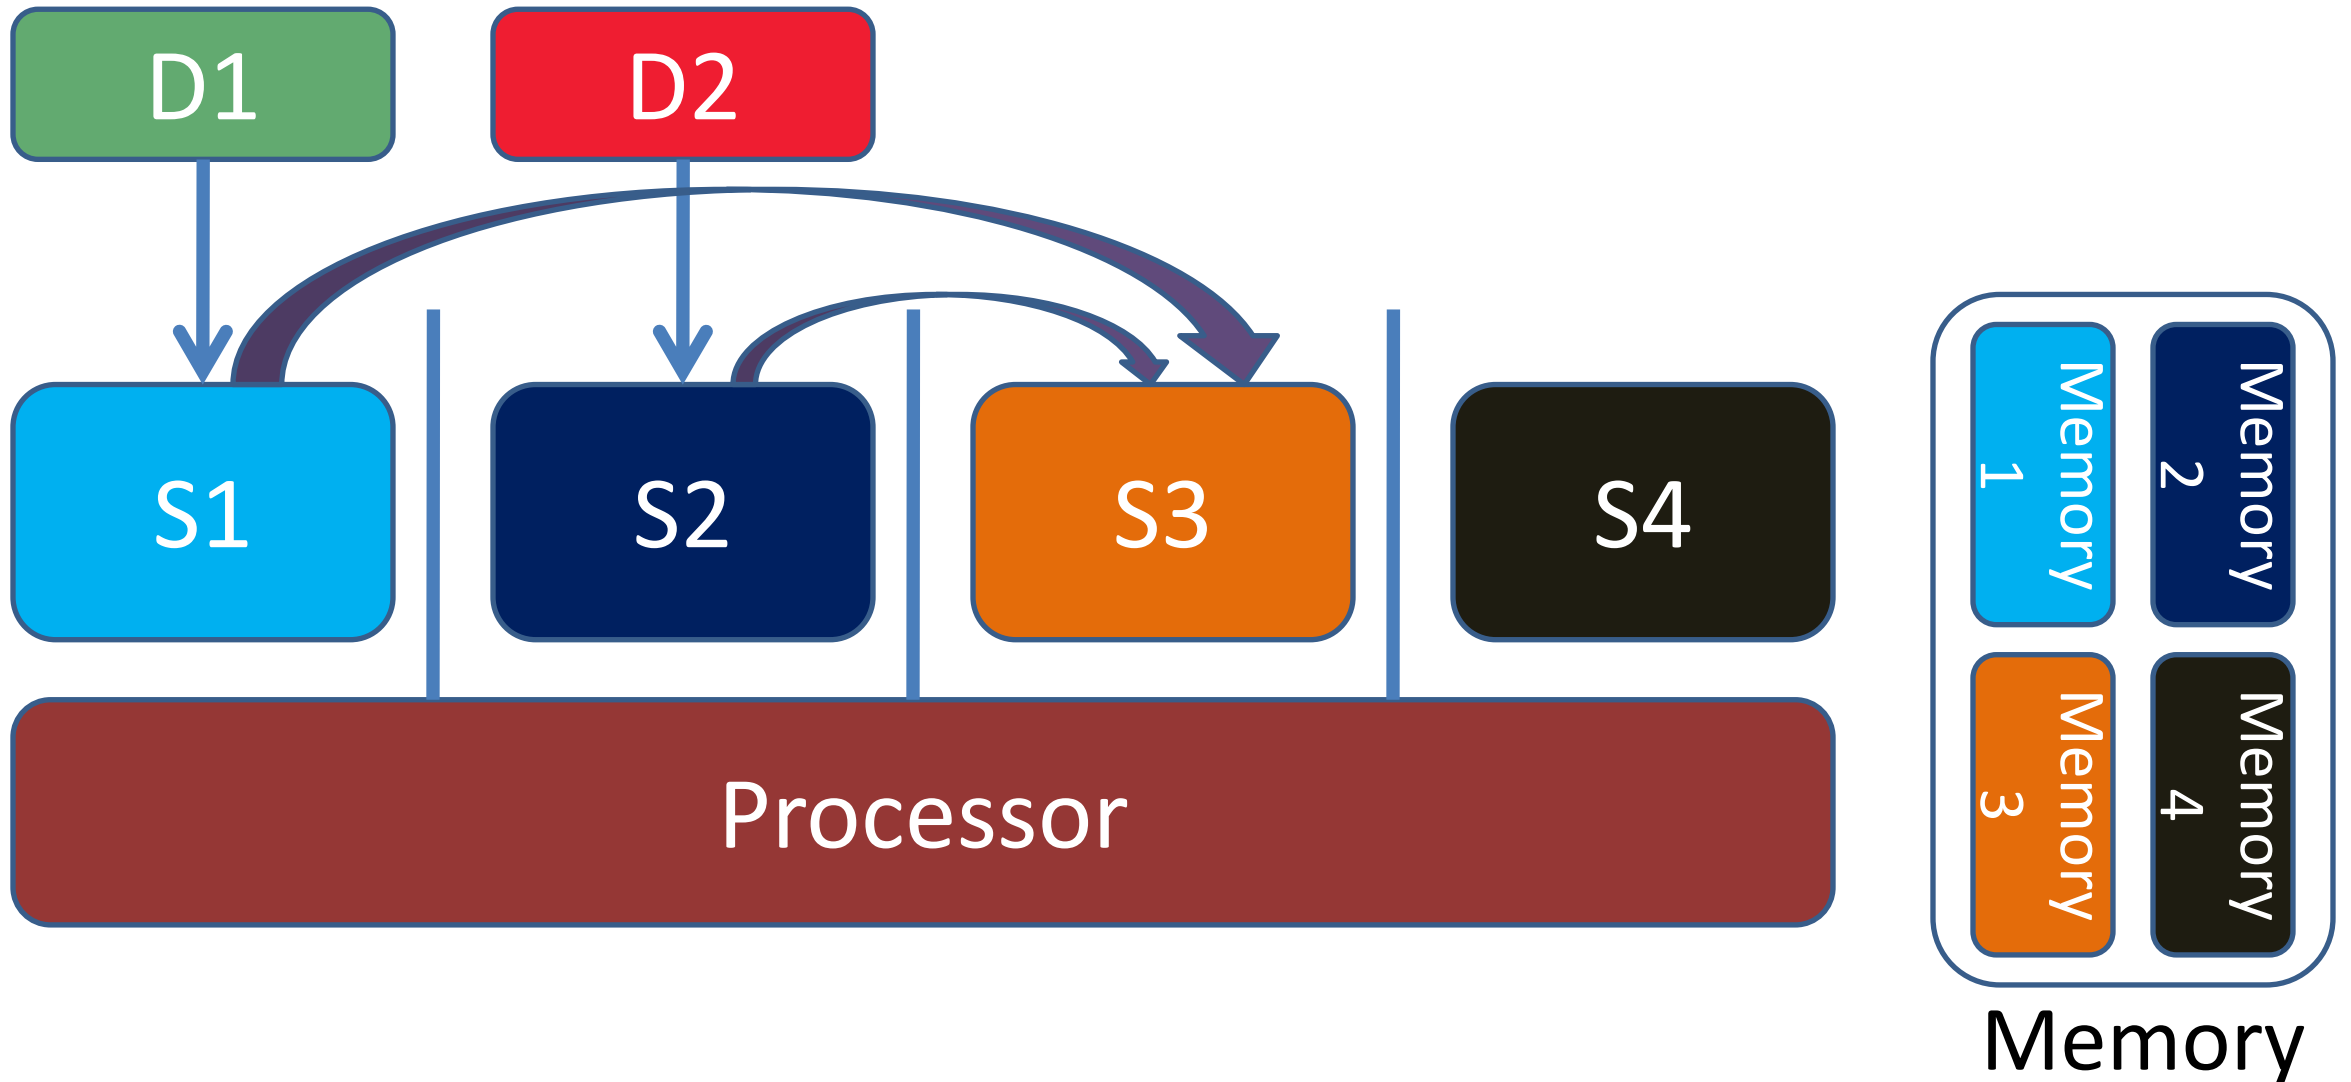
What Muen does?



Example Policy

- 4 subjects, say S1, S2, S3 and S4
- 2 communication channels
 - S1 to S3
 - S2 to S3
- 2 devices D1 and D2 connected to subjects S1 and S2 respectively
- Memory for each subject is of size 512 MB

Example Muen system



Policy

- Contains the following information
 - Memory areas
 - Communication channels
 - Subjects
 - Memory for the subject
 - Devices attached to the subject
 - Channels where it can read or write

Subject

- Abstract view: a full stand-alone machine running a software
- Similar to a virtual machine on a hypervisor
- Can be a bare metal program or an OS
- Also called partition or regime in a separation kernel

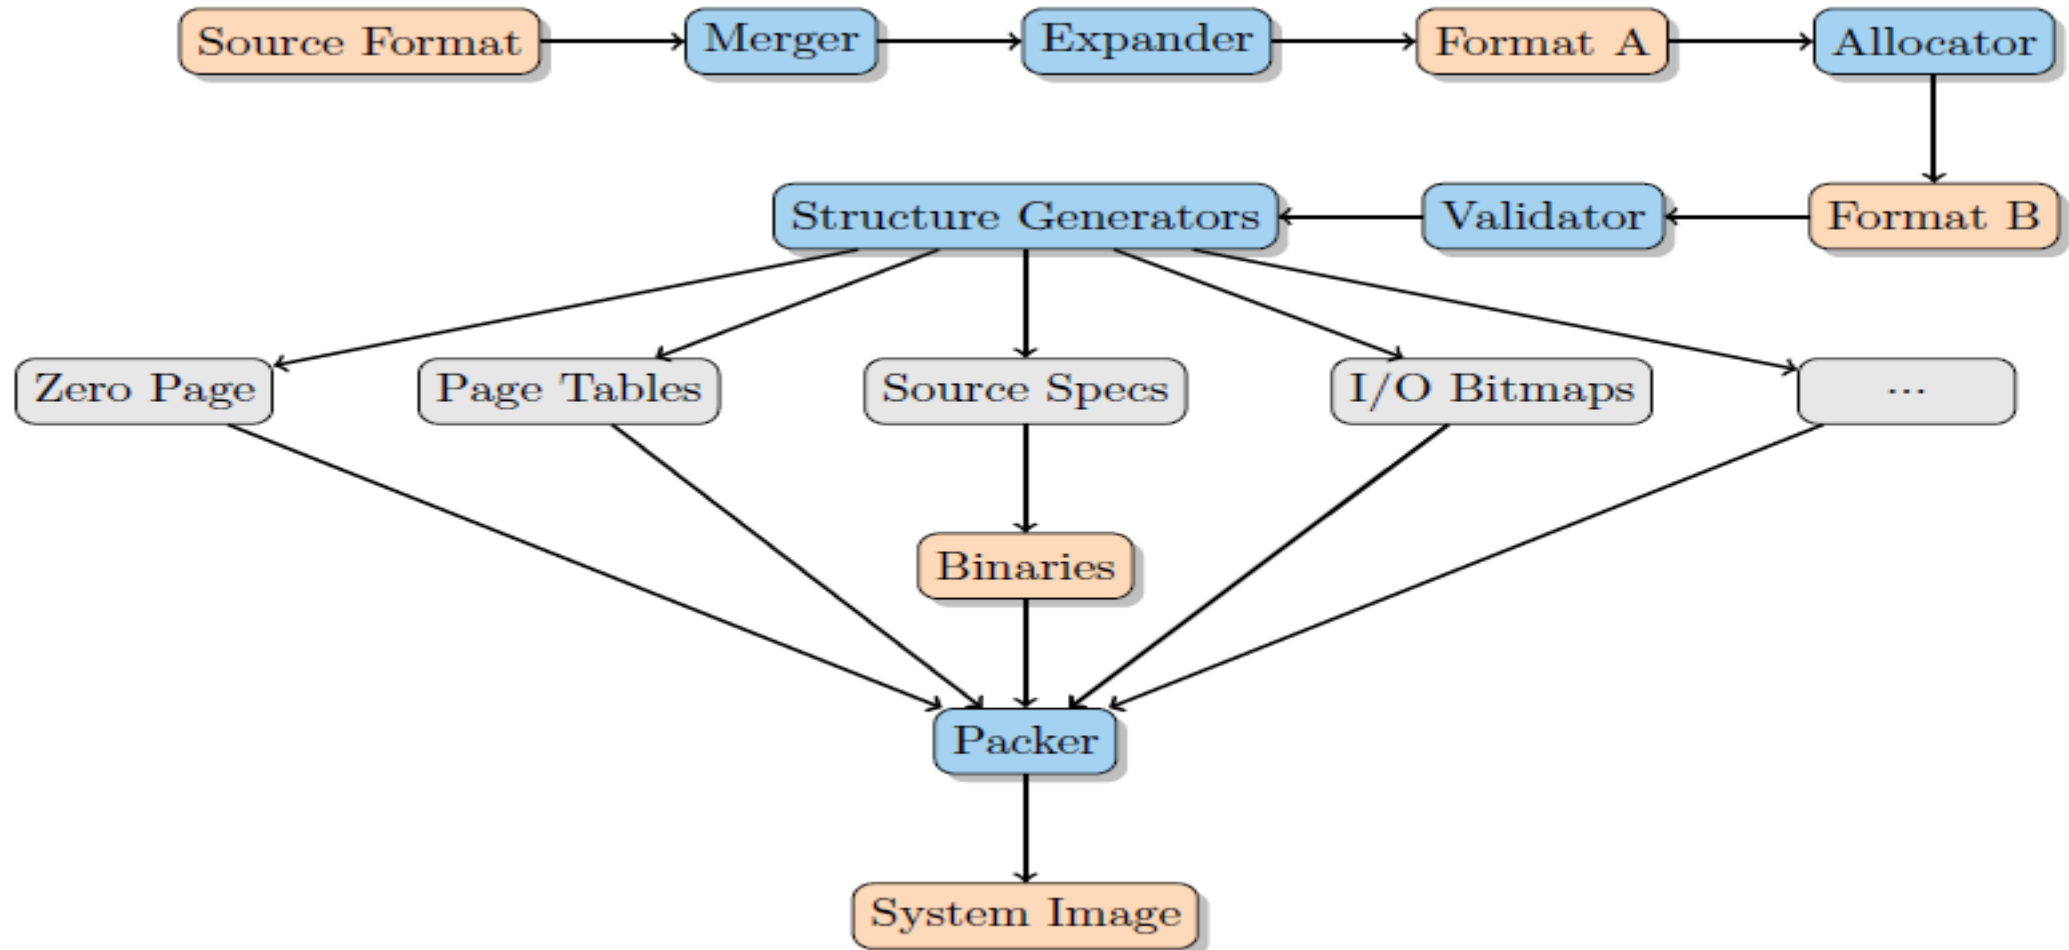
Channels

- Can be either through
 - Shared memory
 - Events
- Complete isolation between subjects except these communication channels
- One-way channels

Scheduling

- Static scheduling
- Round-robin scheduling
- Uses two kinds of frames
 - Minor frames – a subject runs for one minor-frame
 - Major frames – used for synchronization on multiple processors
- 1 Major frame can consist of multiple minor frames

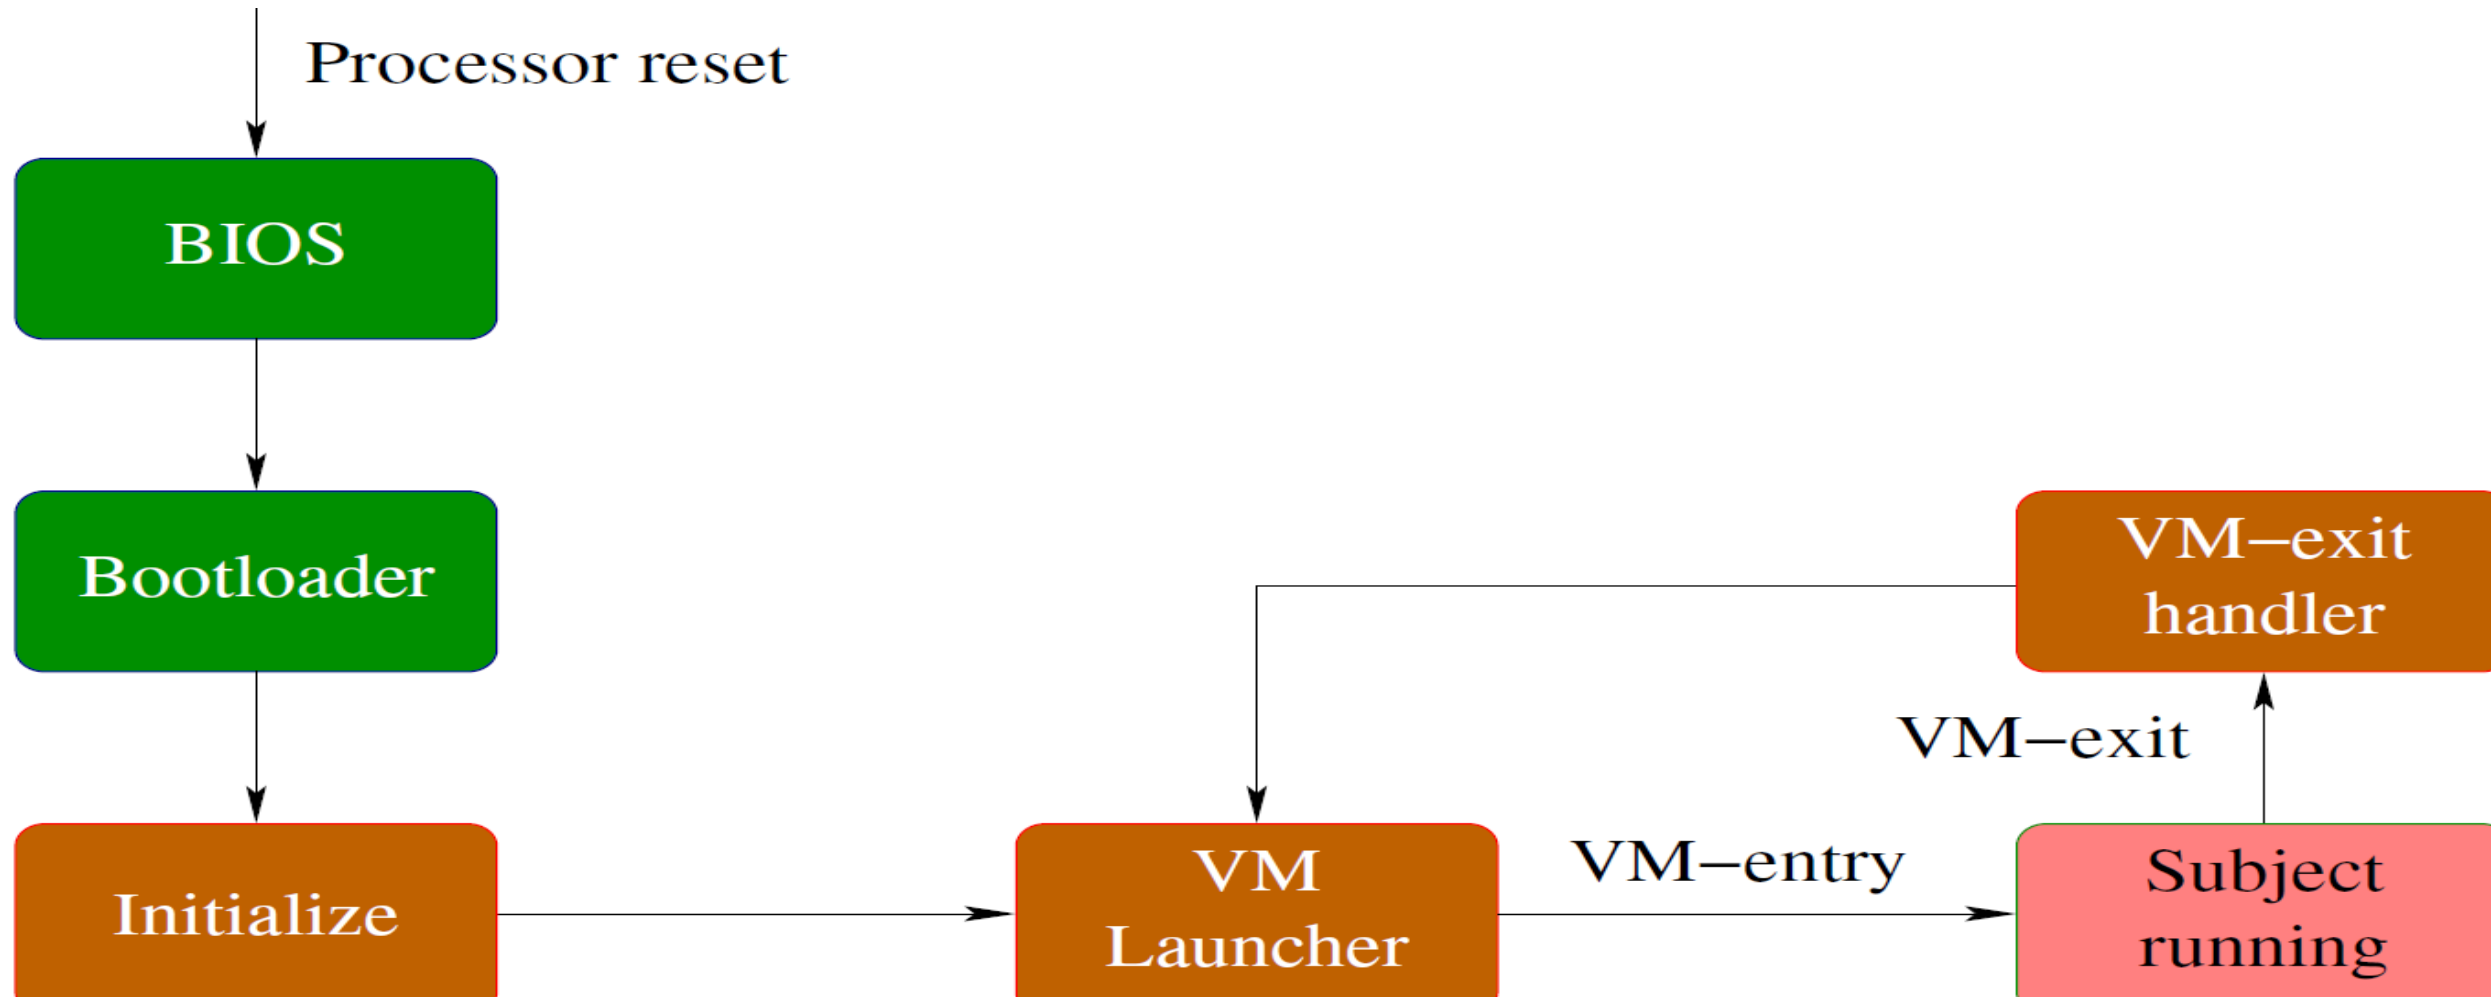
Tool-chain



Source: Muen tool-chain document

Address map

Overview of working of Muen



Initialization

- Set up segmentation and paging
- Set up IDT
- Checking validity of system
- Performing VMXON
- Configure VMCS for each subject
- Initialize scheduler and VMX timer

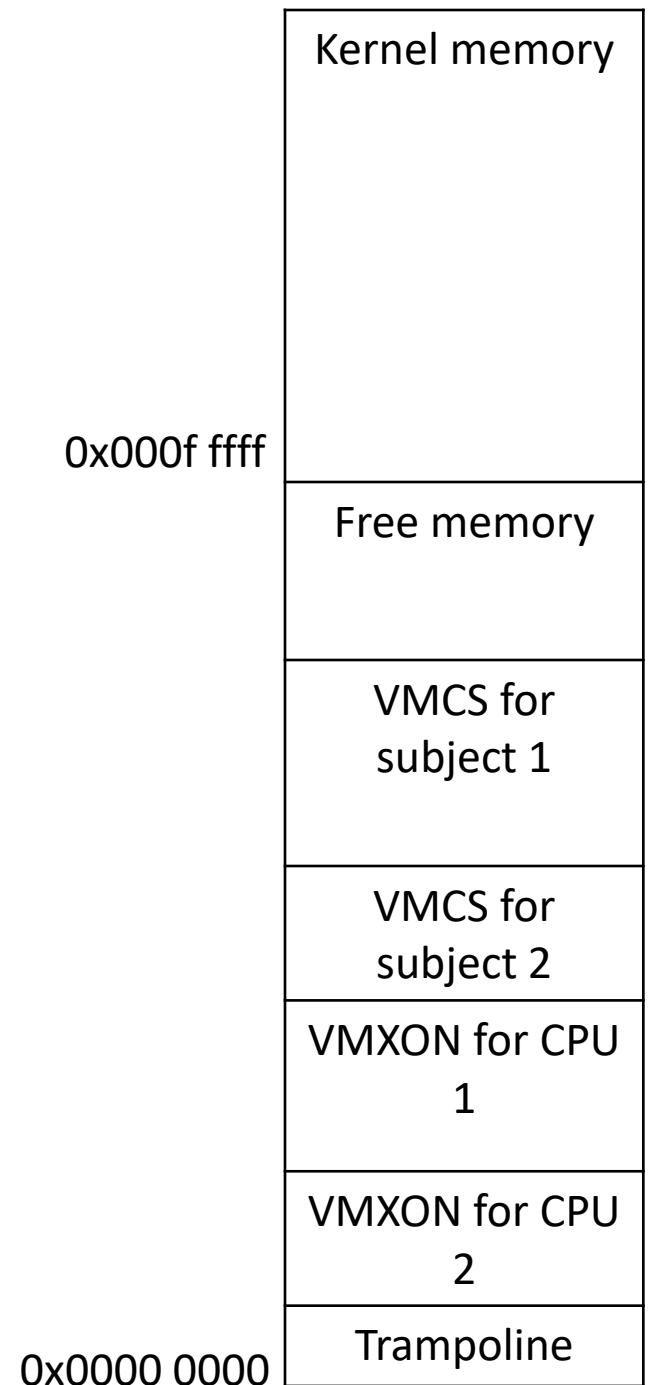
System after initialization



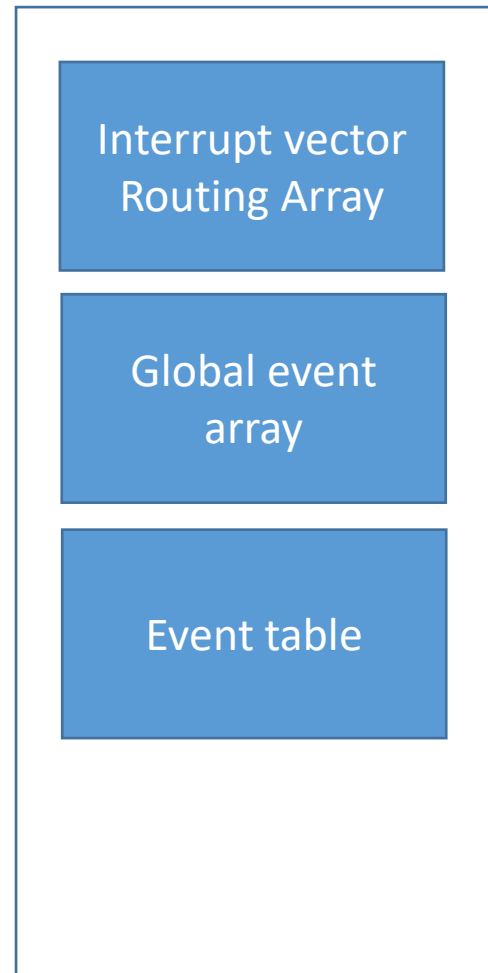
VMCS 1



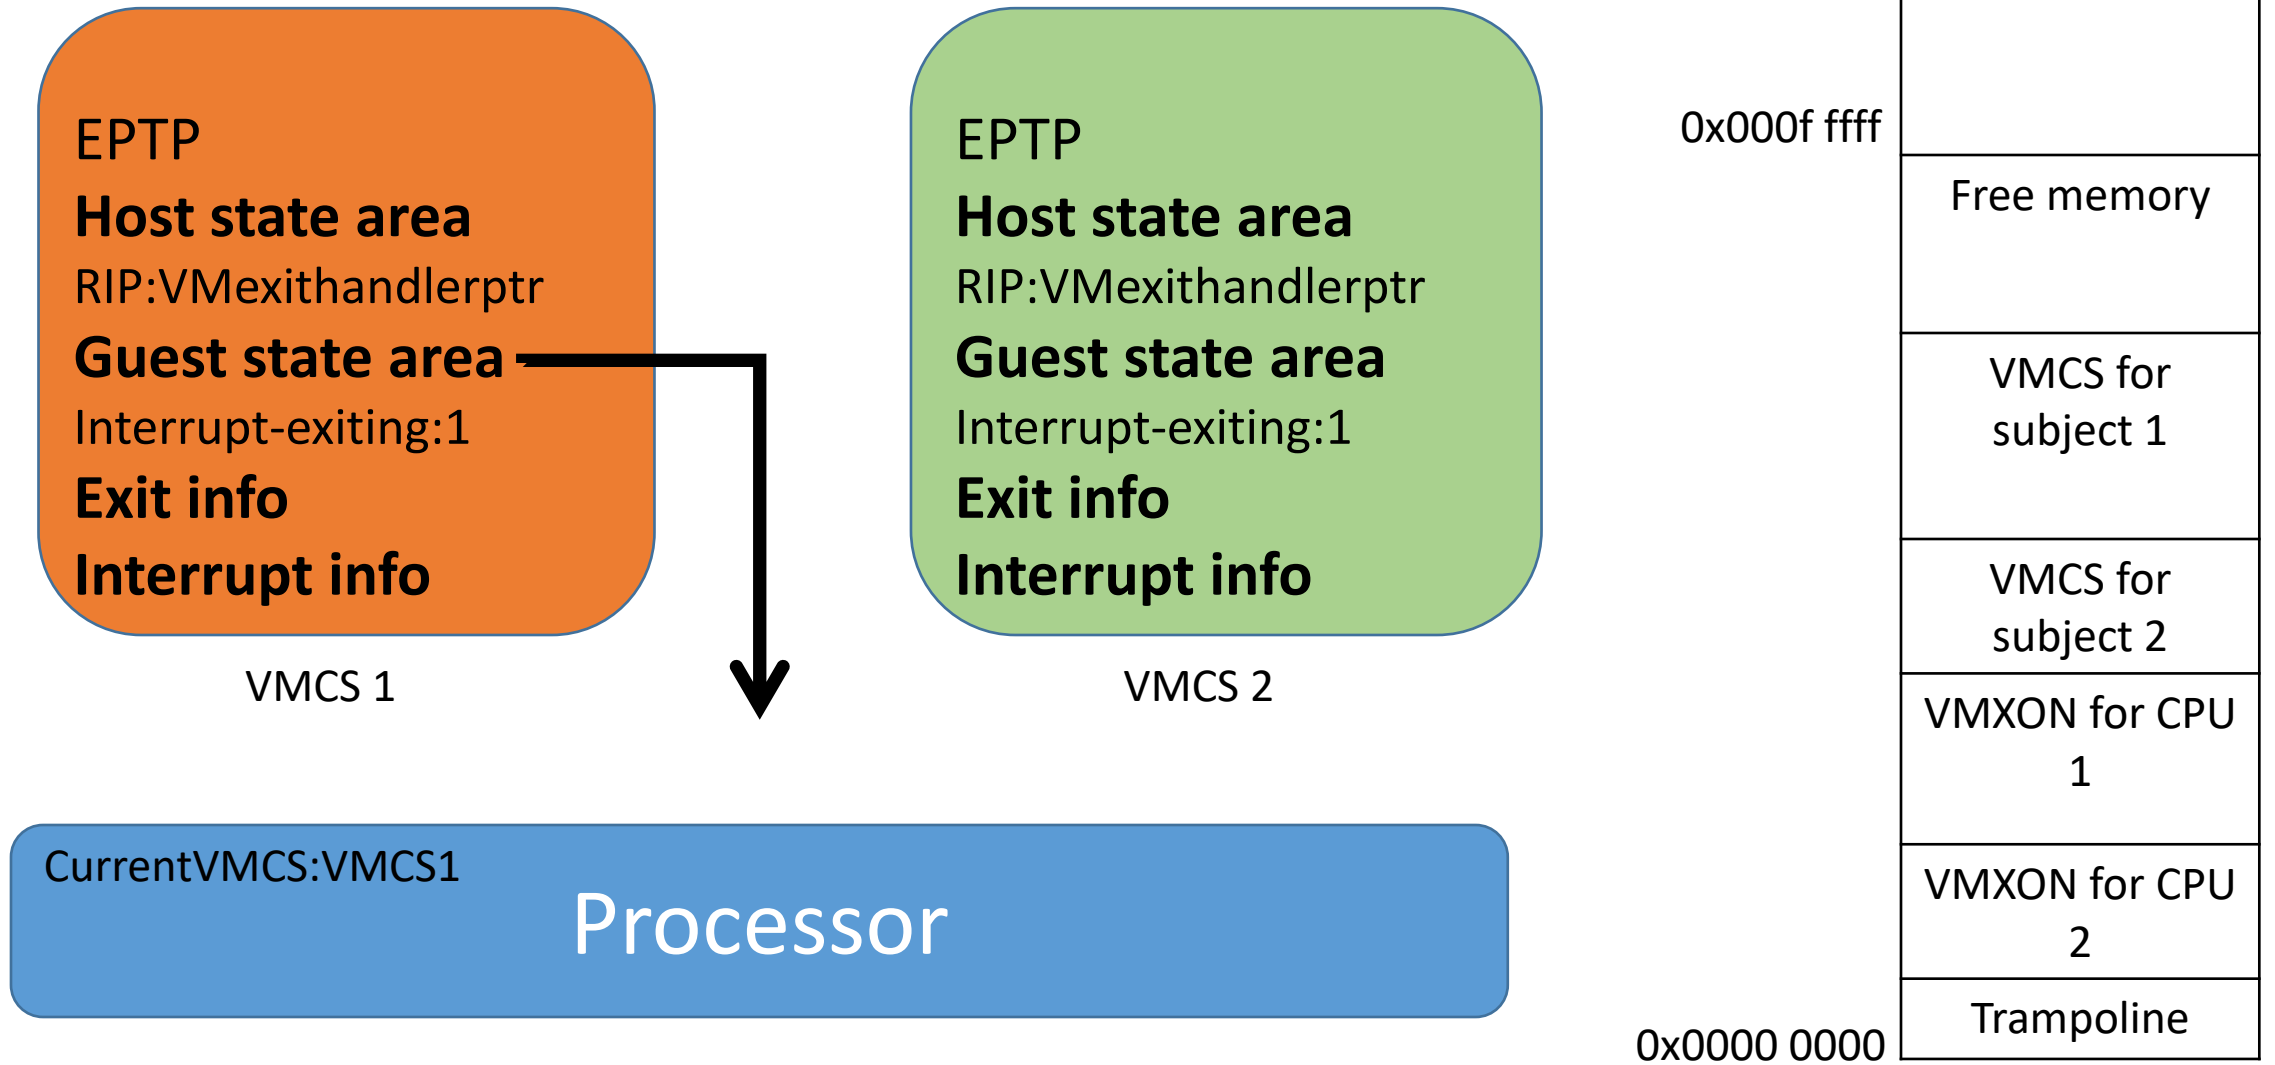
VMCS 2



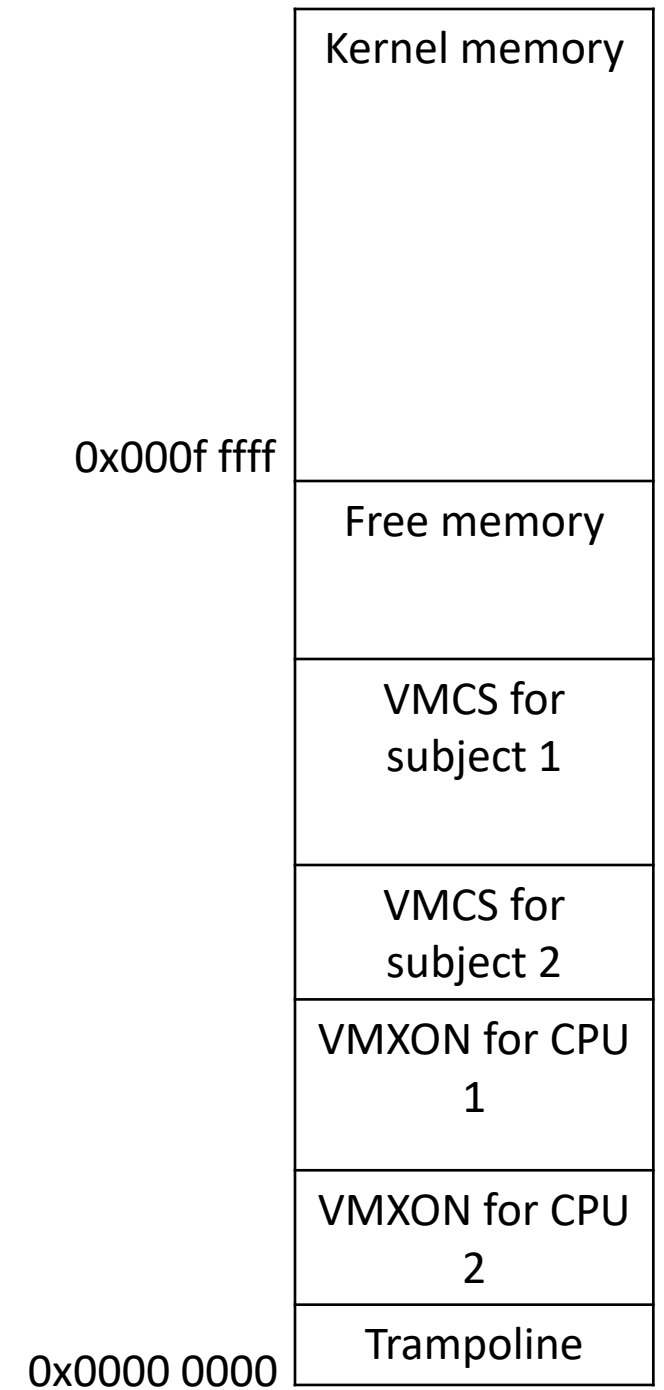
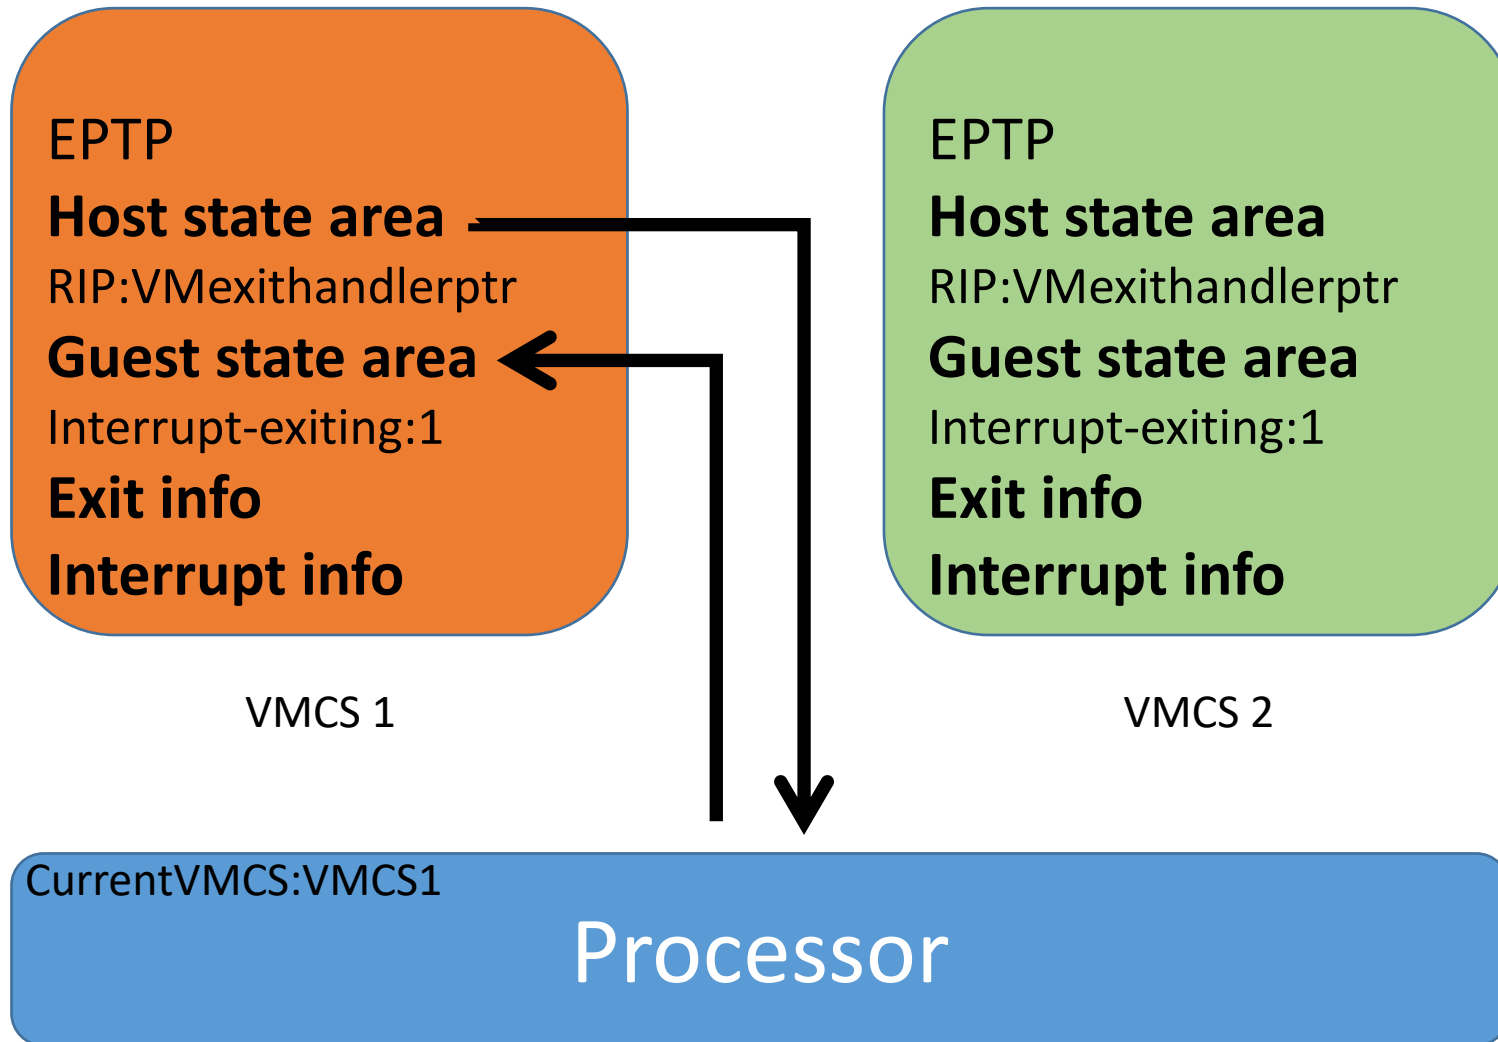
Kernel data structures



VM Entry



VM Exit




VM Exit

- Various reasons of VM-exit
 - External interrupts
 - VMX preemption timer expiry
 - VMCALL instruction
 - Interrupt-window exiting

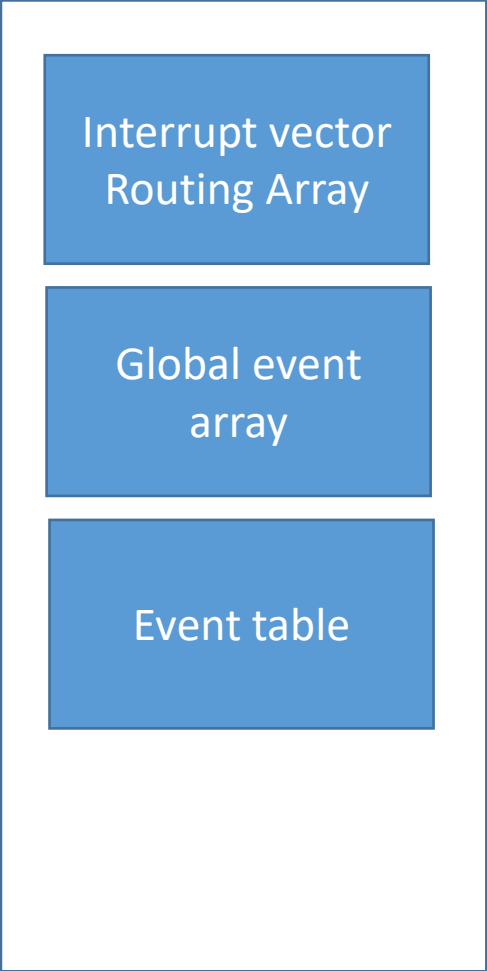
VM Exit – External Interrupt



VMCS 1

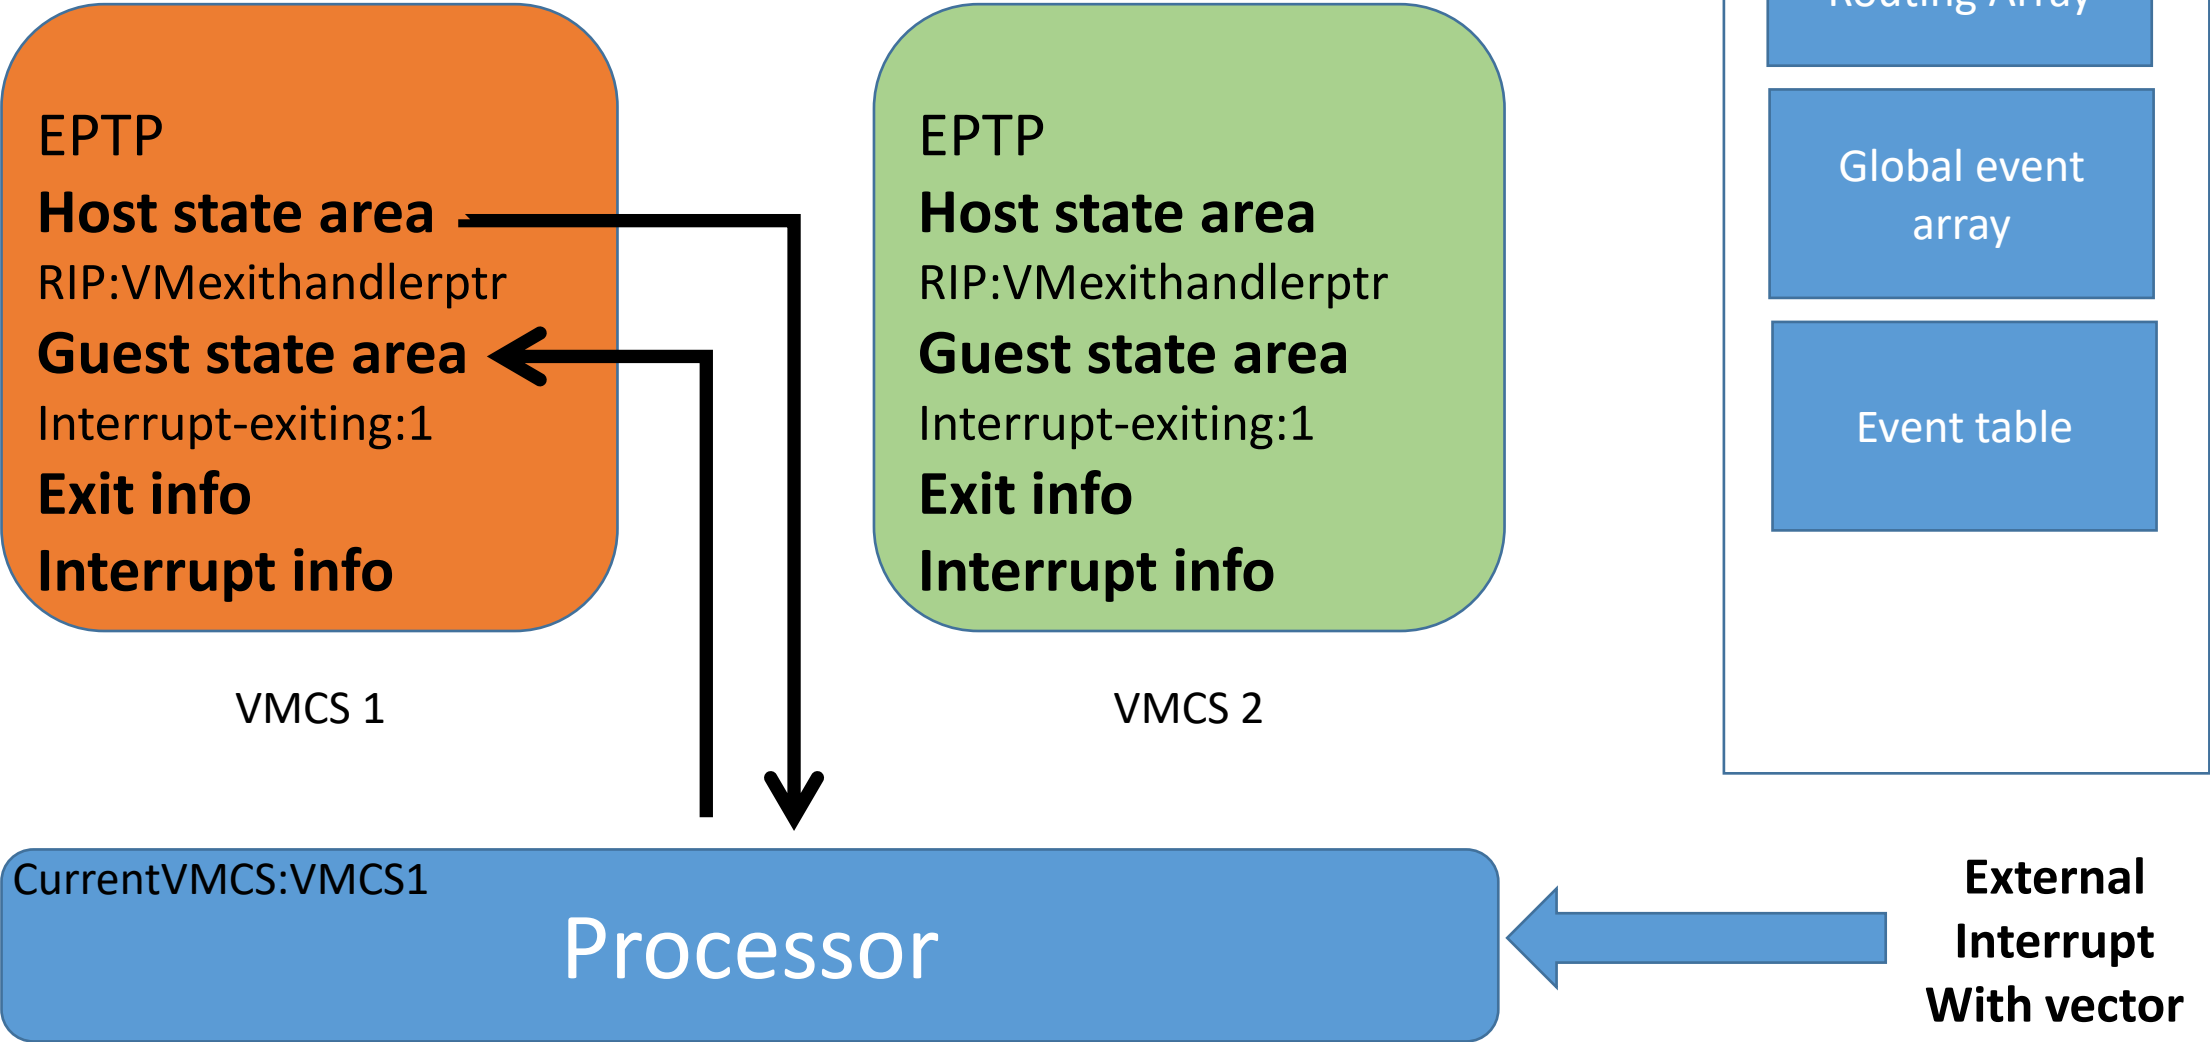


VMCS 2



**External
Interrupt
With vector**


VM Exit – External Interrupt



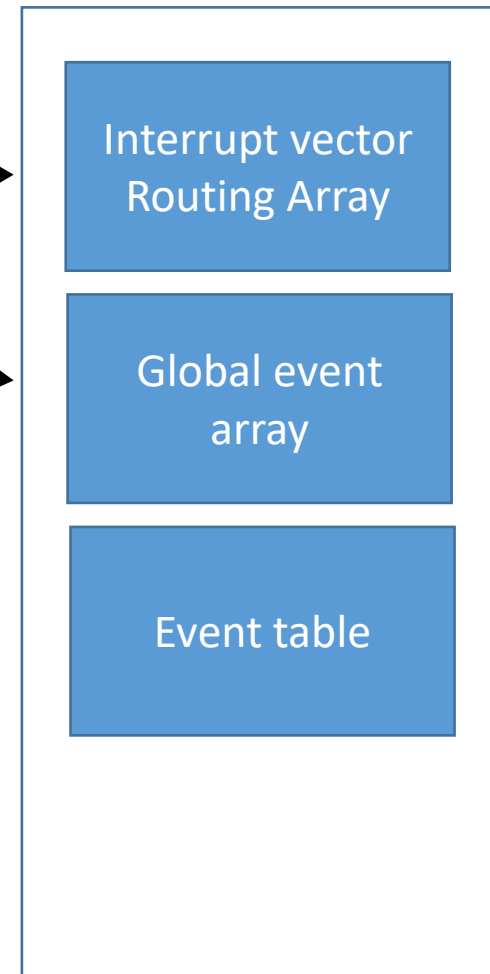
VM Exit – External Interrupt



VMCS 1



VMCS 2




CurrentVMCS:VMCS1

Processor

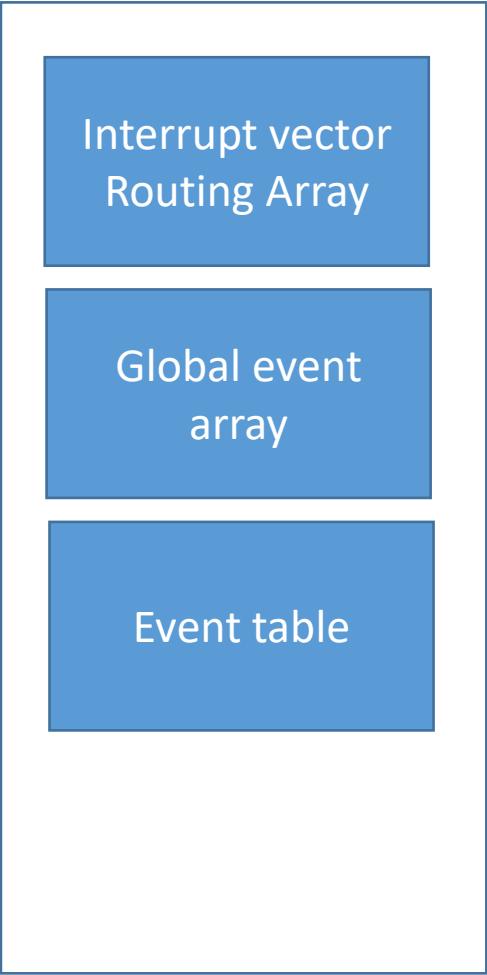
VM Exit – External Interrupt



VMCS 1


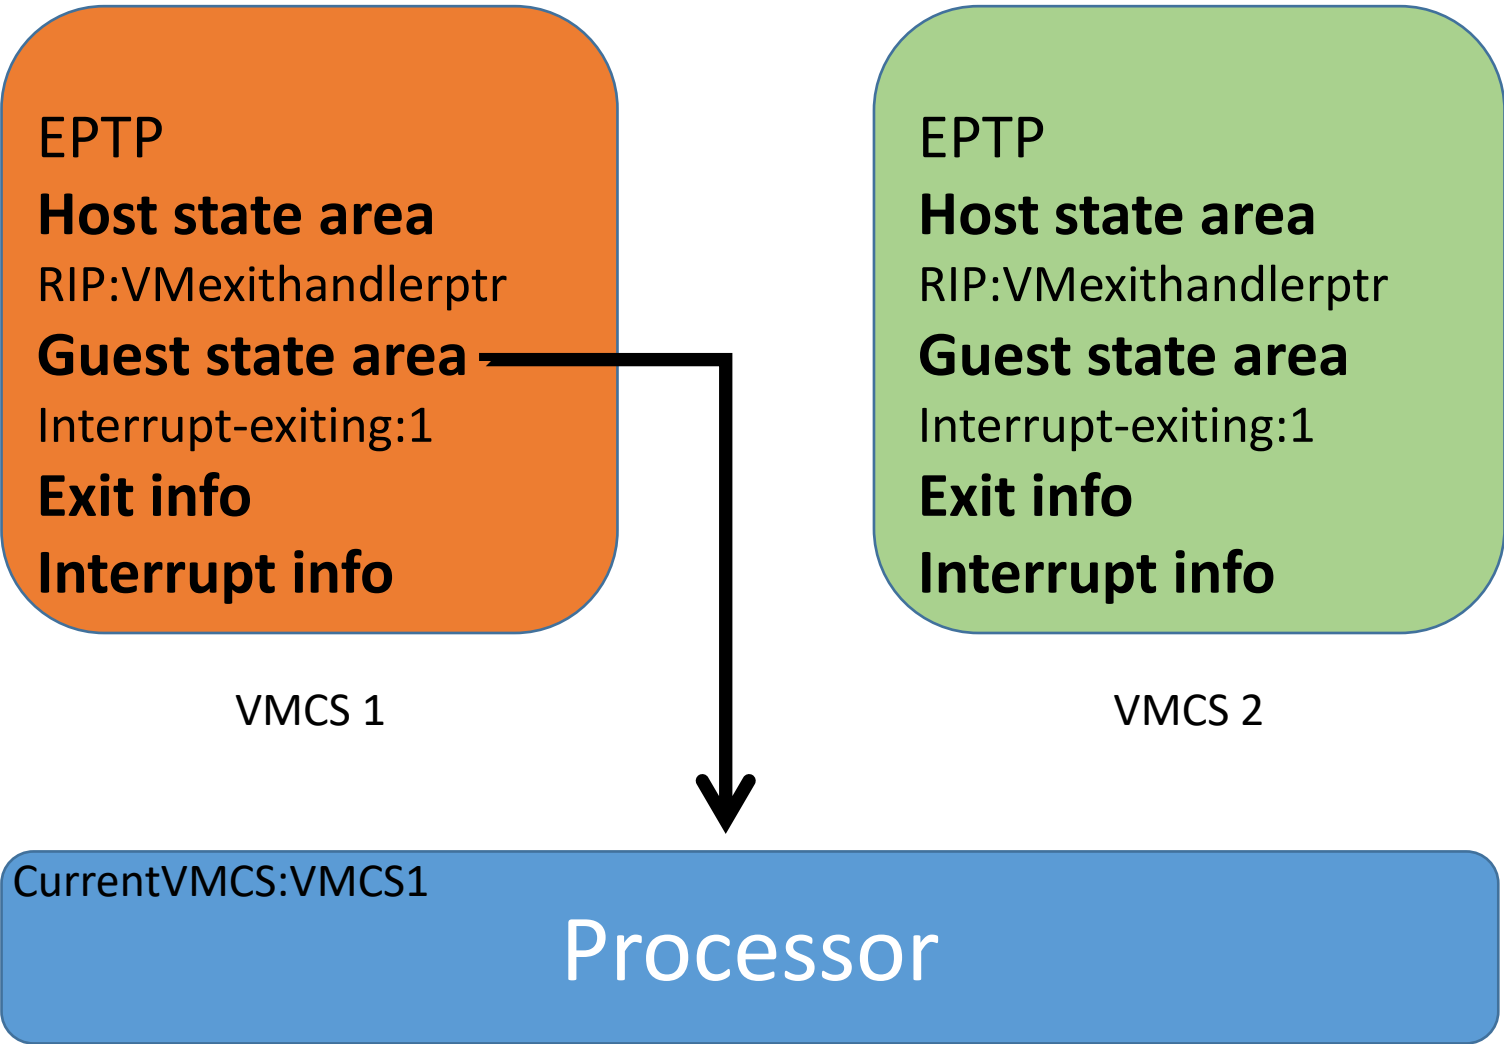


VMCS 2



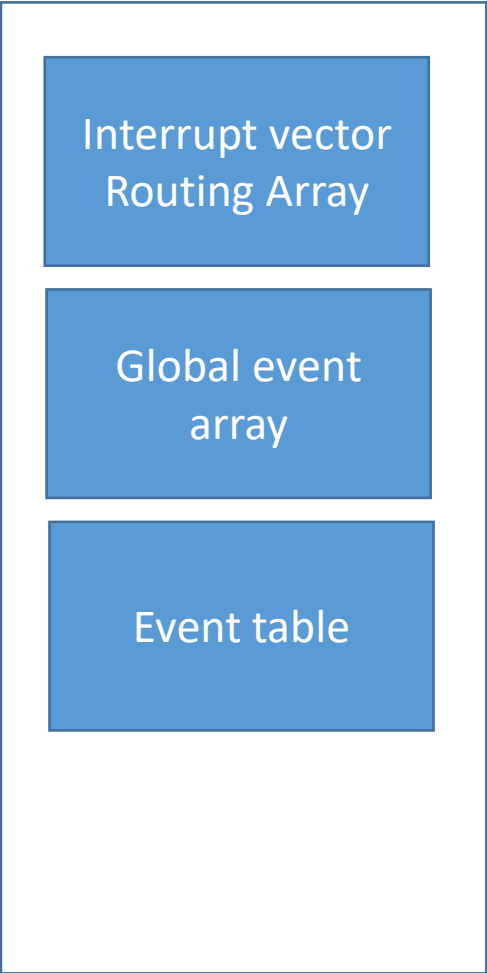
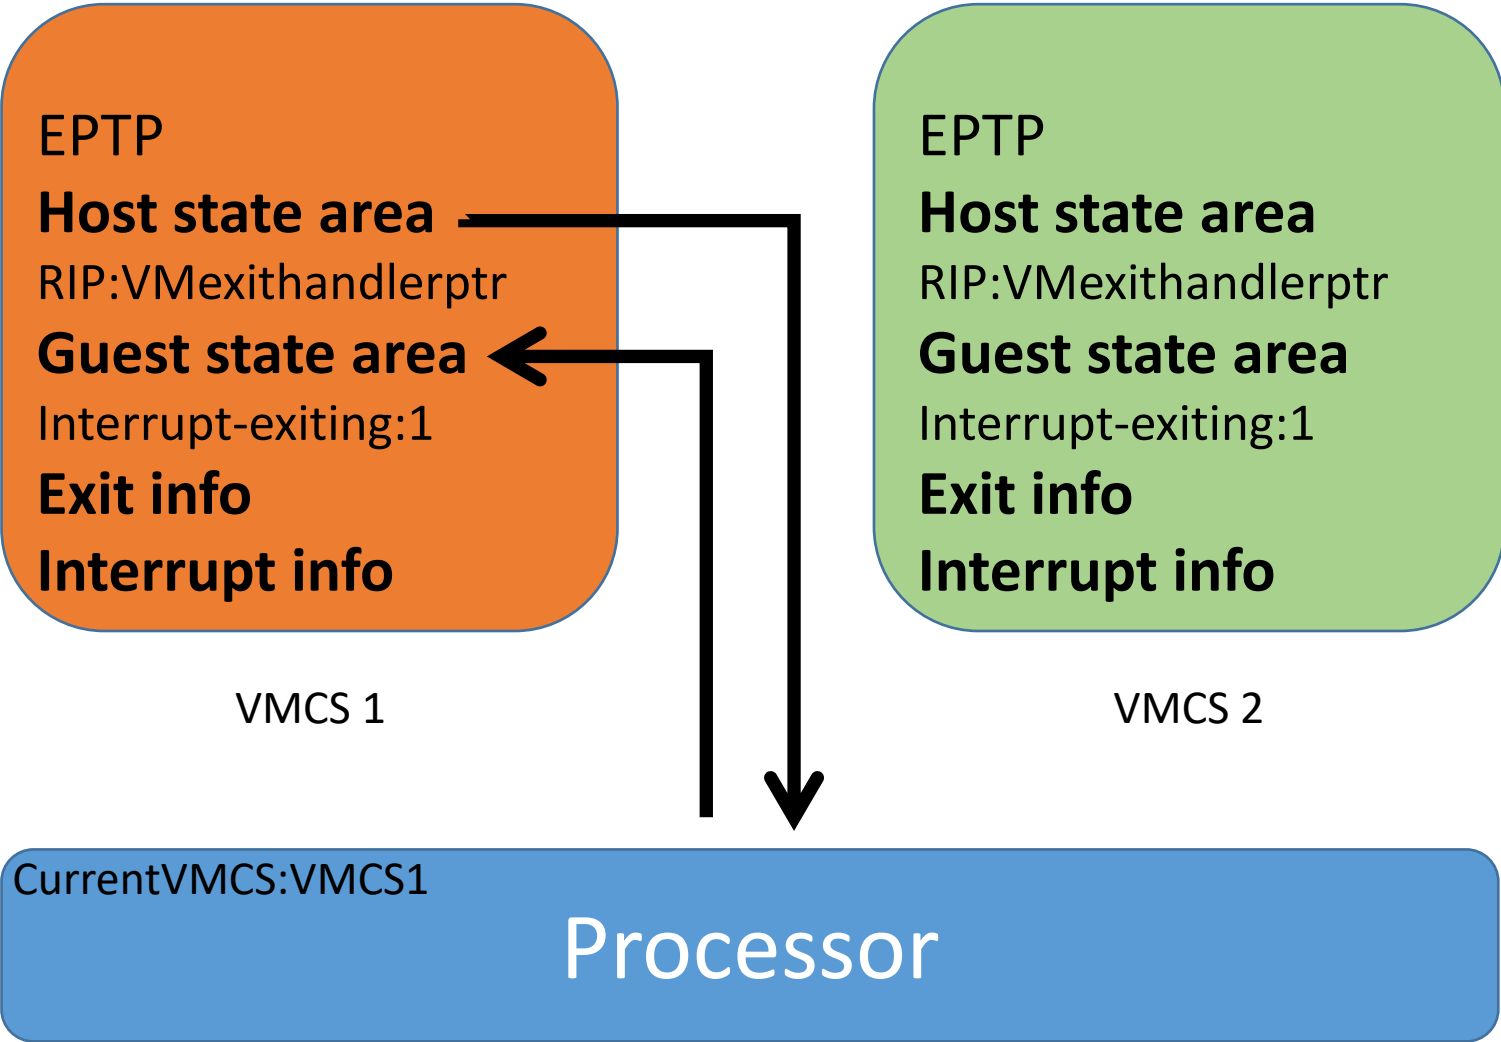
Interrupt handler – subject 2 with vector vn

VM Exit – External Interrupt



Subject 1 starts running again

VM Exit – Timer Expiry



VM Exit – Timer Expiry

EPTP

Host state area

RIP:VMexithandlerptr

Guest state area

Interrupt-exiting:1

Exit info

Interrupt info

VMCS 1

EPTP

Host state area

RIP:VMexithandlerptr

Guest state area

Interrupt-exiting:1

Exit info

Interrupt info

VMCS 2

CurrentVMCS:VMCS2

Processor

Interrupt vector
Routing Array

Global event
array


Event table

VM Exit – Timer Expiry

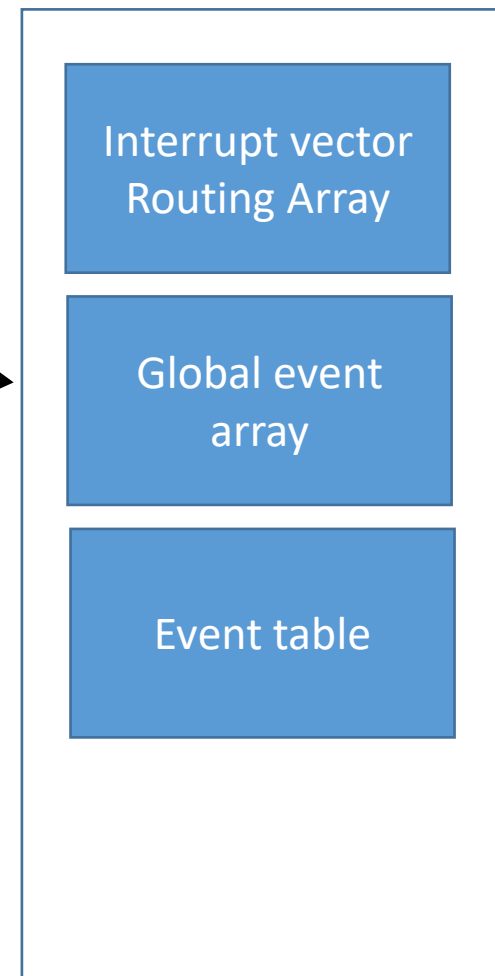
Kernel data structures



VMCS 1



VMCS 2

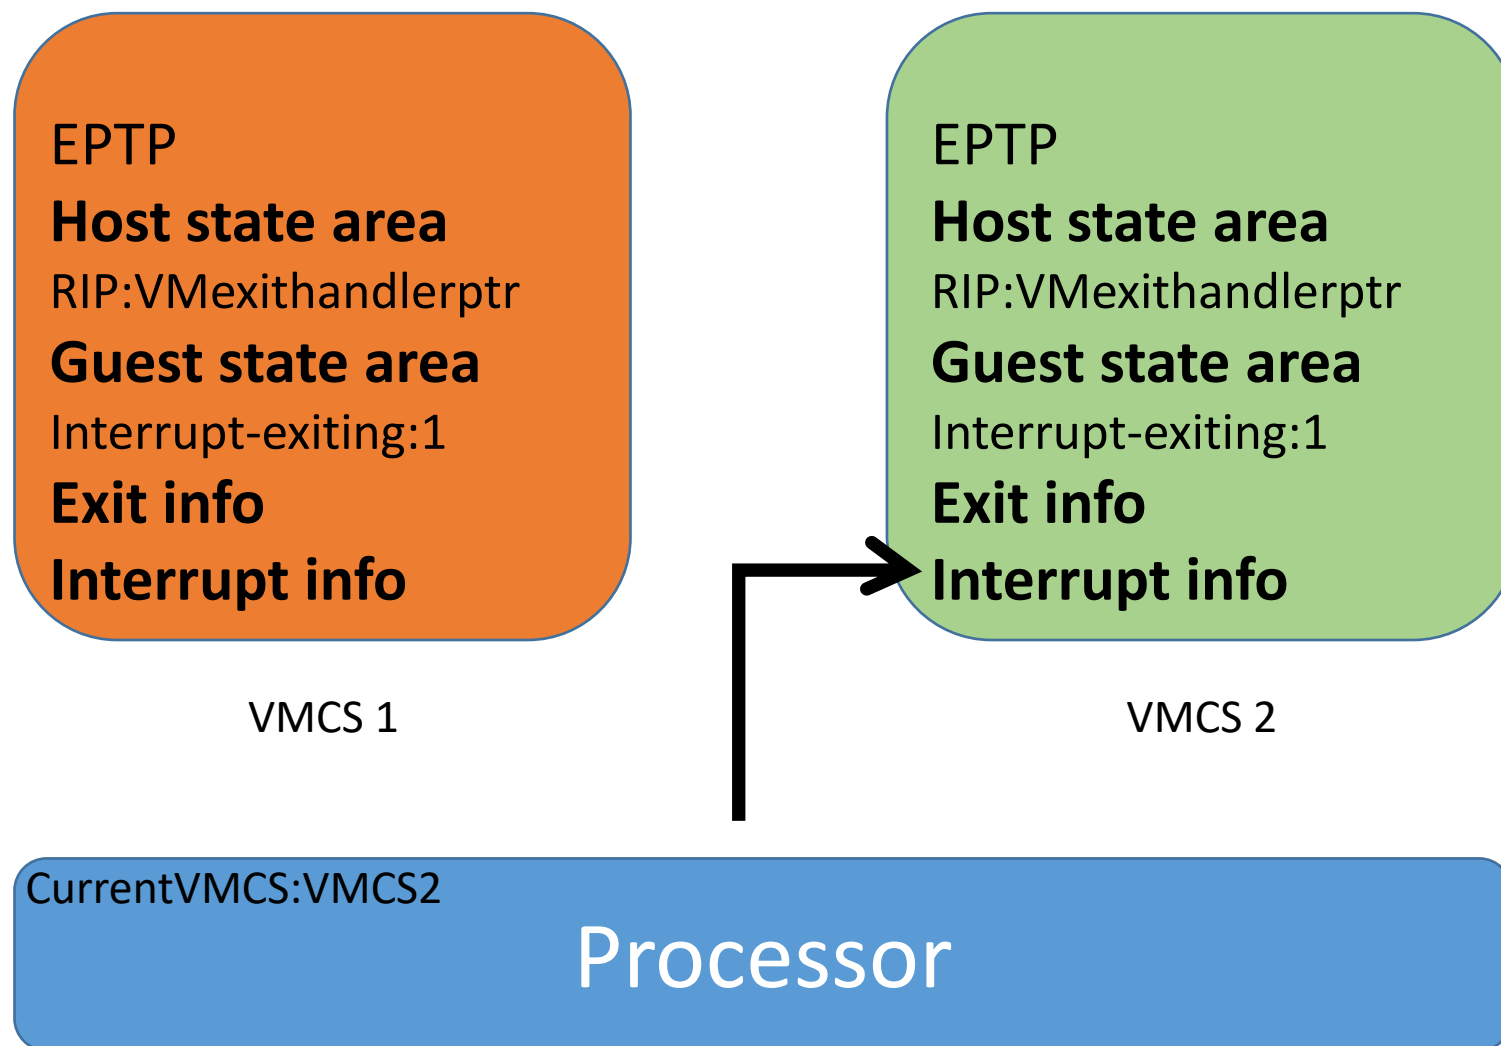


CurrentVMCS:VMCS2

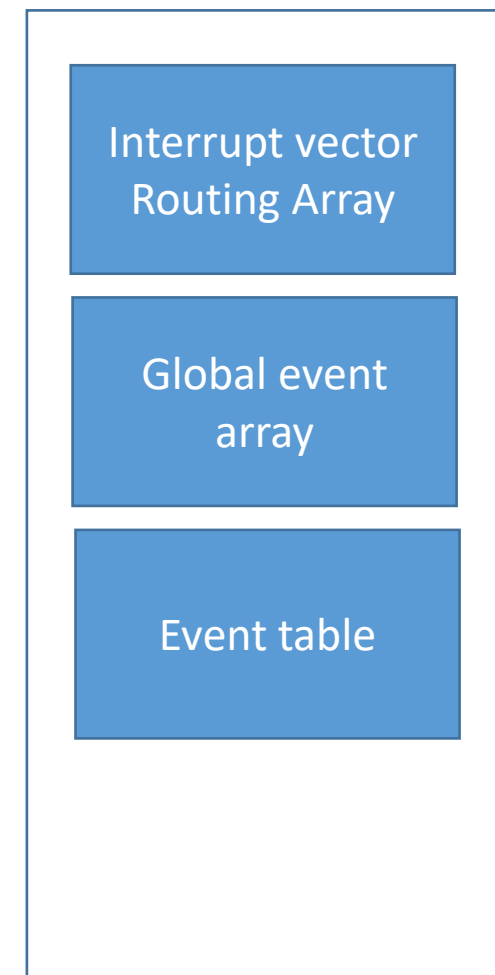
Processor

Checking global event table for subject 2

VM Exit – Timer Expiry

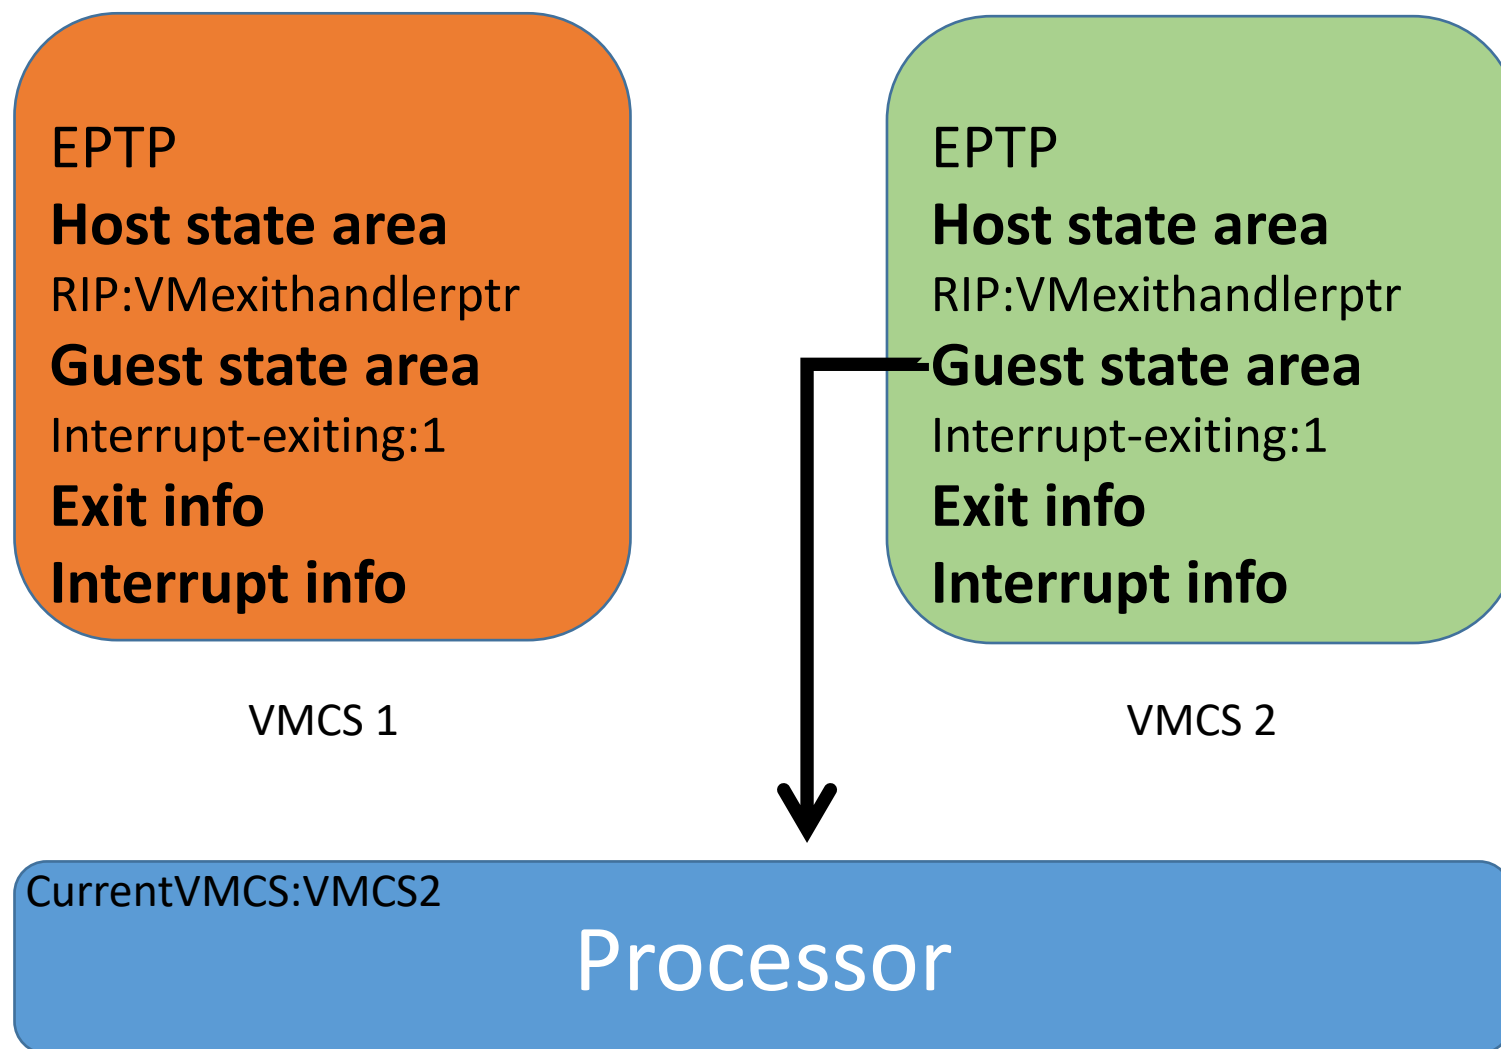


Kernel data structures

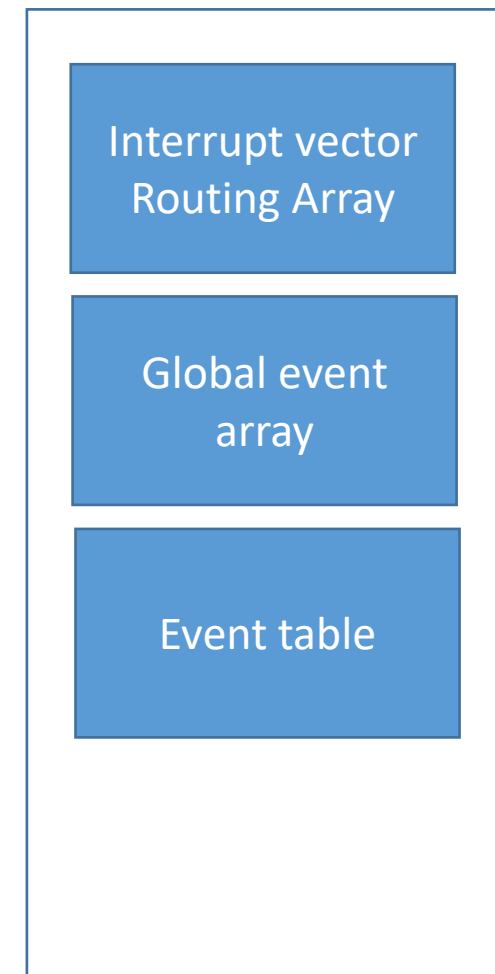


Writing interrupt info
in VMCS of subject 2

VM Exit – Timer Expiry

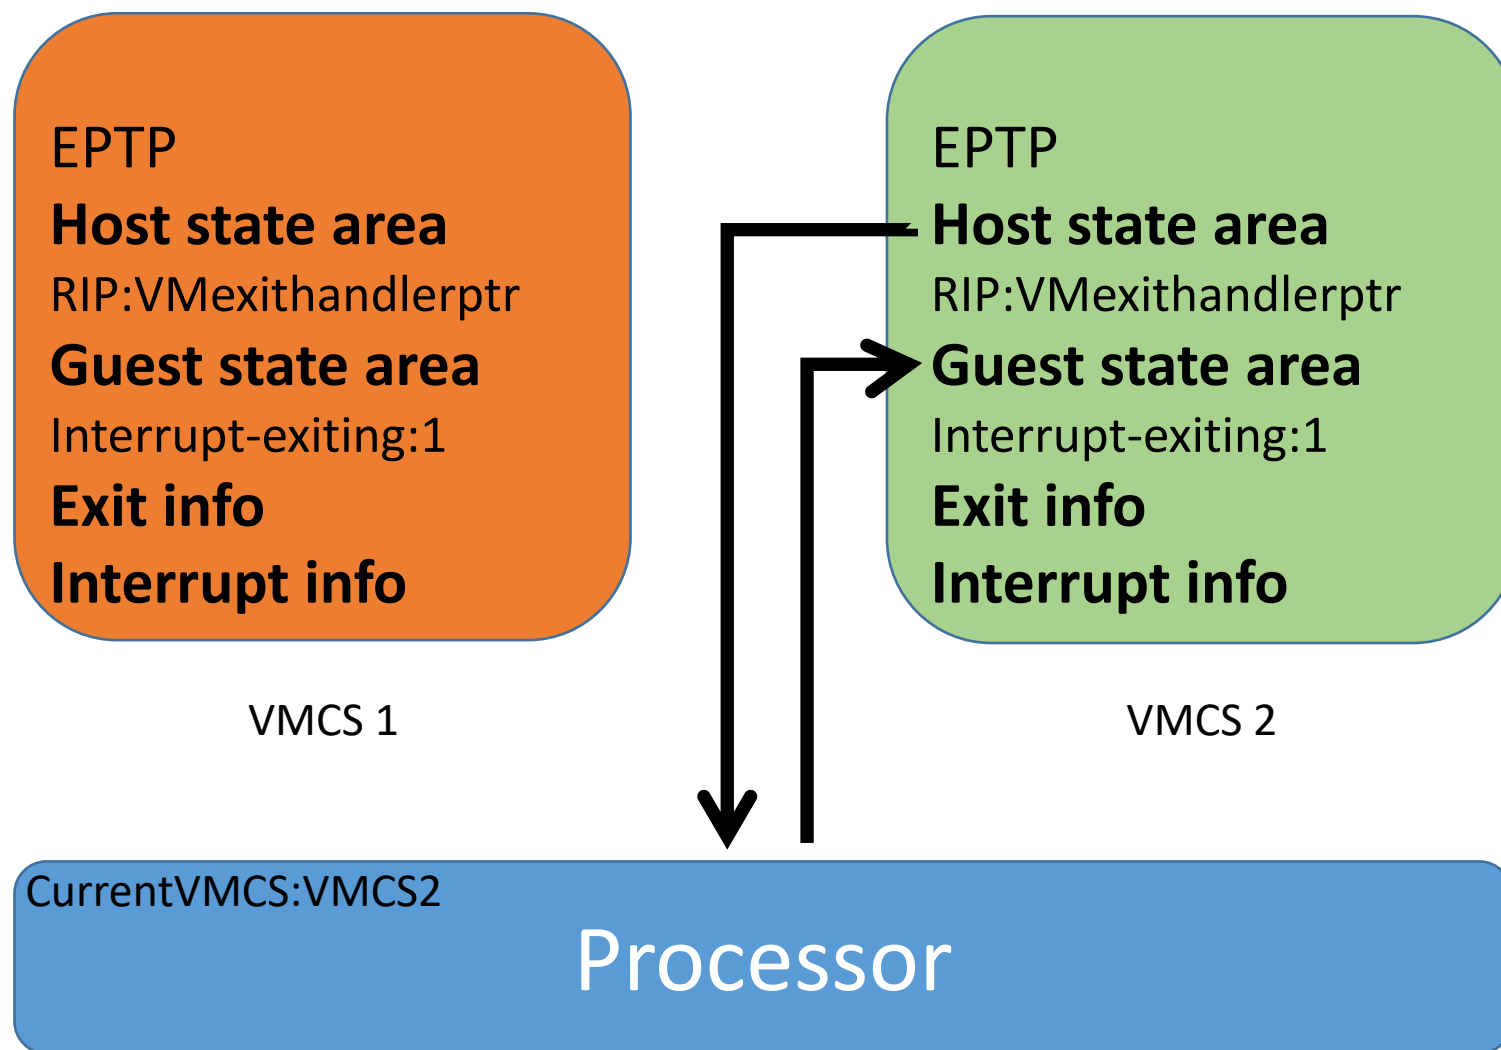


Kernel data structures

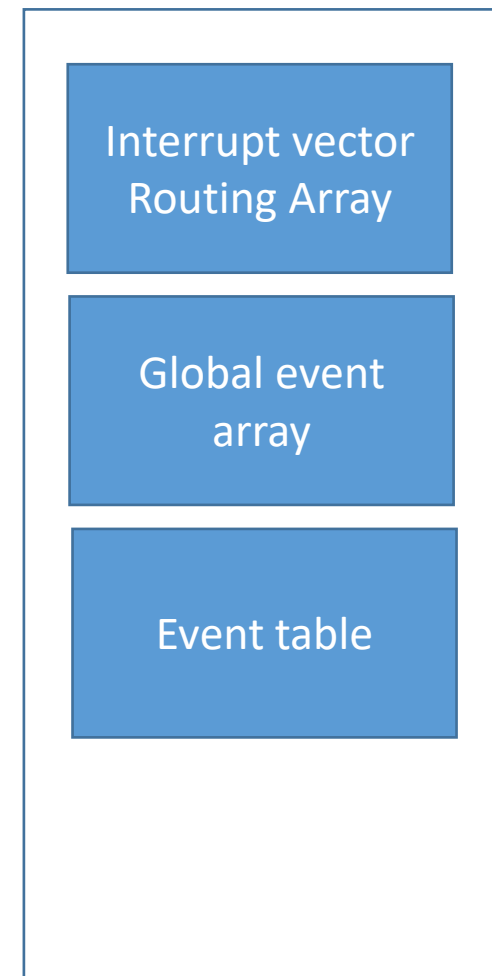


Checking interrupt info field before entry

VM Exit - Hypercall



Kernel data structures




Guest 2 executed VMCALL instruction with operand in A register

VM Exit - Hypercall

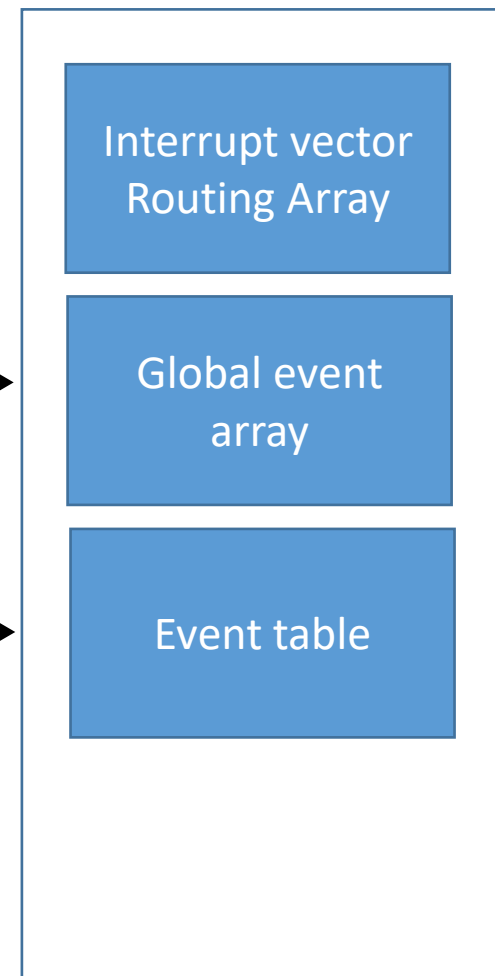
Kernel data structures



VMCS 1



VMCS 2



CurrentVMCS:VMCS2


Processor

Checks event table and set the bit for dest subject with dest event no

VM Exit - Hypercall



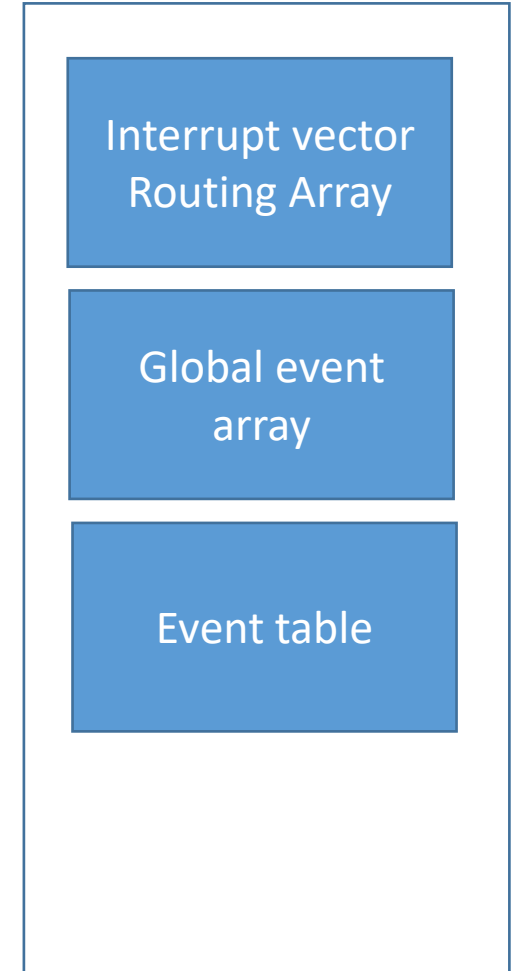
VMCS 1



VMCS 2



Kernel data structures



Starts running again if not
handover event